


BENCHMARK YOUR OWN ENDPOINT

Turn a traffic profile into evidence you can defend.

Replay a measured workload as **open-loop traffic**, score it against the customer's own SLA, preserve the exact run record, and state only what the evidence supports.

The decision path

Each gate can stop the run. No green status is inferred from HTTP 200 alone.



What one valid run can establish

- Caller-experienced latency and delivery at the exact tested arrival schedule.
- Whether caller-supplied TTFT, full-generation, and answer-rate targets held.
- HTTP status, stream completion, parse integrity, usage coverage, and observed 429 evidence.
- Achieved cache reuse only where the endpoint reports cached-token usage.

What it does not establish

- A fixed-rate pass is not an endpoint ceiling, provider headroom, or future capacity guarantee.
- Network-path timing is context only. No network estimate is subtracted from TTFT or full-generation latency.
- Structural answer validity is not semantic correctness or task quality.
- Do not quote provider billing or full-run cost without independently validated rates and complete attempt-level usage.

The evidence contract

Keep the complete directory. A browser or PDF view is navigation, not the source seal.

start.json Exact configuration, source state, input identities, and declared experiment.	requests.jsonl Recorded metadata rows for setup and replay request phases.	summary.json Canonical metrics and independent decision dimensions.	manifest.json Schema-v3 artifact hashes, byte counts, row counts, and run identity.	report.html / .md Readable views derived from the same canonical summary.
--	--	---	---	---

STAGE ONE

Prove mechanics with minimal, quota-planned traffic.

This shipped profile is deliberately illustrative. It tests transport, parsing, timing, usage, and answer completeness. **It is not Decagon demand and cannot support a performance or capacity claim.**

Authorization

Confirm the workspace, endpoint, operator, stop conditions, and approved test window.

Authentication

Use an origin-matched profile on the standard workspace route. This resolver does not mint route-optimized scoped tokens.

Local gate

Run the full test suite and local oracle before contacting the endpoint.

LOCAL INSTRUMENT GATE

NO PROVIDER TRAFFIC

```
python3 -m pytest
python3 -m traffic_replay validate --port 0 --format json
```

Current Enterprise P2T quota snapshot

Provider fact dated 2026-08-07; operator rechecked 2026-08-08;
freshness window 7 days.

200,000

INPUT TOKENS / MINUTE

20,000OUTPUT TOKENS / MINUTE;
MAX_TOKENS RESERVED**7,200**

QUERIES / HOUR

80%

HARNESS WARNING BOUNDARY



Quota evidence is intentionally narrow. The gate budgets offered max_tokens, worst-case physical POST attempts, preflight, calibration, and replay against the dated snapshot. It cannot see unrelated workspace traffic, cannot verify the asserted workspace tier, and does not prove provider headroom. Recheck the live source before every paid run. Do not add a service-tier override to this command.

GLM 5.2 INSTRUMENT CANARY

REAL, BILLABLE REQUESTS

```
python3 -m traffic_replay benchmark \
  --host https://YOUR-WORKSPACE-HOST \
  --endpoint databricks-glm-5-2 \
  --auth-profile YOUR-DATABRICKS-PROFILE \
  --profile configs/profile_glm52_canary_illustrative.json \
  --fixed-rate 0.25 --duration 2 \
  --rate-limits configs/rate_limits_databricks_glm_5_2_enterprise_p2t_2026-08-07.json \
  --out-dir results/glm52-instrument-canary \
  --label "INSTRUMENT CONFORMANCE ONLY - NOT CUSTOMER DEMAND OR CAPACITY"
```

1. Exact representatives

Profile mode sends its concrete p50 and p95 requests. Prompt mode sends the first two approved prompts, cycling only if one exists.

2. Exact pass condition

HTTP 200 plus visible content or a valid tool call, clean stream completion, and zero parse errors.

3. Exact refusal

If either representative is incomplete, the measured load stops unless the operator explicitly uses --force.



Managed GLM 5.2 reasoning boundary: Databricks currently classifies databricks-glm-5-2 as reasoning-only and publishes no GLM-specific value that disables it. The tool ships no off control and guesses none. Only explicitly supplied --probe-extra-body candidates are tested, one real request each. An answered probe still does not prove the provider applied the candidate or disabled reasoning; rerun the complete two-representative preflight and inspect reasoning evidence.

[Databricks limits and quotas](#)[Databricks reasoning-model guide](#)

STAGE TWO

Run the customer contract, then separate SLA from capacity.

Build configs/profile_measured.json from approved request metrics, review its digest and dropped rows, then set rates, targets, and client bounds from the authorized experiment.

MEASURED FIXED-RATE RUN

TEMPLATE - REPLACE EVERY YOUR_* VALUE

```
python3 -m traffic_replay benchmark \
  --host https://YOUR-WORKSPACE-HOST \
  --endpoint YOUR-ENDPOINT-NAME \
  --auth-profile YOUR-DATABRICKS-PROFILE \
  --profile configs/profile_measured.json \
  --fixed-rate YOUR-AUTHORIZED-RPS --duration 300 \
  --max-concurrency YOUR-TESTED-WORKER-BOUND \
  --max-pending-requests YOUR-TESTED-PENDING-BOUND \
  --ttft-p95 YOUR-TTFT-MS --ttfg-p95 YOUR-TTFG-MS \
  --success-rate YOUR-FRACTION-0.1 --rate-limits RATE_LIMITS.json \
  --fail-on caution --out-dir results/customer-fixed-rate
```

R **Arrival rate is the load axis.** Do not use legacy `--concurrency`. `--max-concurrency` is only a client safety bound; the report shows the in-flight concurrency that the offered rate actually carried. TTFT is scored with the selected caller definition. Full-generation latency is reported separately.

GUARDED RATE SWEEP

SEQUENTIAL AND STATEFUL

```
python3 -m traffic_replay sweep \
  --host https://YOUR-WORKSPACE-HOST --endpoint YOUR-ENDPOINT-NAME \
  --auth-profile YOUR-DATABRICKS-PROFILE \
  --profile configs/profile_measured.json \
  --rate YOUR-RUNG-1,YOUR-RUNG-2,YOUR-RUNG-3 \
  --duration 120 --cooldown 60 \
  --cpt YOUR-PREMEASURED-CHARACTERS-PER-TOKEN \
  --max-concurrency YOUR-TESTED-WORKER-BOUND \
  --max-pending-requests YOUR-TESTED-PENDING-BOUND \
  --ttft-p95 YOUR-TTFT-MS --ttfg-p95 YOUR-TTFG-MS \
  --success-rate YOUR-FRACTION-0.1 --rate-limits RATE_LIMITS.json \
  --out-dir results/customer-rate-sweep
```

! Cooldown is spacing, not proof that rolling quotas, autoscaling, cache state, or unrelated traffic reset. The whole requested ladder is budgeted before preflight. The built-in snapshot gate covers direct Databricks pay-per-token routes only; other modes need their own guard. Stop on 429s, parse errors, answer failures, target misses, load shortfall, pending drops, generator saturation, or production impact. The highest valid rung tested is still not an endpoint ceiling.

Read five decisions, not one green badge

A target pass can coexist with invalid measurement evidence. Preserve the qualification.

DIMENSION	QUESTION ANSWERED	REQUIRED READING
Evidence integrity	Do current bytes match the manifest and completion chain?	Verify externally; a fresh run cannot verify the manifest that encloses it.
Measurement validity	Did errors, delivery, parsing, sample size, or missing evidence undermine inference?	CAUTION and INVALID qualify every downstream number.
Acceptance checks	Did the run meet the customer's explicit targets?	PASS, MISS, INCONCLUSIVE, or NOT EVALUATED.
Quota state	Was an HTTP 429 observed in any recorded request phase?	Zero observed 429s is not quota headroom.
Endpoint capacity	Did the endpoint hold at this exact tested load?	Never translate one held rung into a ceiling or guarantee.

Create a separate run-verification receipt

The verifier re-reads the sealed source, cross-checks the canonical journal and summary, and writes a new receipt without modifying the run.

```
python3 -m traffic_replay verify-run \
  results/customer-fixed-rate/RUN-DIRECTORY \
  --out results/receipts/customer-fixed-rate

python3 -m traffic_replay verify-sweep \
  results/customer-rate-sweep
```

Receipt boundary

verification.json, verified-report.html, and verified-report.md are sealed in the receipt. SHA-256 consistency is not a digital signature: it does not prove authorship, trusted time, repository availability, or immutability after verification.

Before sharing: review labels, endpoint metadata, file paths, and configuration. Run artifacts are designed to retain metadata and hashes rather than raw request or response content; the measured extra_body remains part of reproducibility evidence.

github.com/debu-sinha/llm-traffic-replay