


BENCHMARK YOUR OWN ENDPOINT

Turn a traffic profile into evidence you can defend.

Replay a measured workload as **open-loop traffic**, score it against customer-owned acceptance targets, preserve the exact run record, and state only what the evidence supports.

The decision path

Each gate can stop the run. No green status is inferred from HTTP 200 alone.



What one valid run can establish

- Caller-experienced latency and delivery at the exact tested arrival schedule.
- Whether customer-owned TTFT, full-generation, and answer-rate targets held.
- HTTP status, stream completion, parse integrity, usage coverage, and observed 429 evidence.
- Achieved cache reuse only where the endpoint reports cached-token usage, plus response-model identity and pre-run/post-drain endpoint stability.

What it does not establish

- A fixed-rate pass is not an endpoint ceiling, provider headroom, or future capacity guarantee.
- Network-path timing is context only. No network estimate is subtracted from TTFT or full-generation latency.
- Structural answer validity is not semantic correctness or task quality.
- Do not quote provider billing or full-run cost without independently validated rates and complete attempt-level usage.

The evidence contract

Keep the complete directory. A browser or PDF view is navigation, not the source seal.

start.json Exact configuration, source state, input identities, and declared experiment.	requests.jsonl One metadata row per request operation, with per-attempt admission events, attempt counts, and the terminal outcome.	summary.json Canonical metrics and exactly five independent decision dimensions.	manifest.json Schema-v3 artifact hashes, byte counts, row counts, and run identity.	report.html / .md Readable views derived from the same canonical summary.
--	---	--	---	---

STAGE ONE

Prove mechanics with minimal, quota-planned traffic.

This shipped profile is deliberately illustrative. It tests transport, parsing, timing, usage, and answer completeness. **It is not measured customer demand and cannot support a performance or capacity claim.**

Authorization

Confirm the workspace, endpoint, operator, stop conditions, and approved test window.

Authentication

Use an origin-matched profile on the standard workspace route. This resolver does not mint route-optimized scoped tokens.

Local gate

Run the full test suite and local oracle before contacting the endpoint.

LOCAL INSTRUMENT GATE

NO PROVIDER TRAFFIC

```
python3 -m pytest
python3 -m traffic_replay validate --port 0 --format json
```

Published Enterprise P2T defaults: tier and headroom not verified

Dated 2026-08-07; rechecked 2026-08-08. Verify Enterprise tier and shared-workspace headroom before paid traffic.

200,000 INPUT TOKENS / MINUTE	20,000 OUTPUT TOKENS / MINUTE; MAX_TOKENS RESERVED	7,200 QUERIES / HOUR
200 QUERIES / SECOND / WORKSPACE	4,000,000 SERIALIZED REQUEST BYTES; CONSERVATIVE 4 MB BOUNDARY	80% HARNESS WARNING BOUNDARY

- !

Quota evidence is intentionally narrow. A single command-local, no-wait guard atomically reserves every physical inference POST immediately before conn.request. It enforces QPS, QPH, input/output TPM, and serialized request bytes across preflight, probes, calibration, and replay. A local denial sends no POST and is not an HTTP 429 or endpoint-capacity result. The guard cannot see unrelated workspace traffic, verify the asserted tier, or prove provider headroom. Recheck the live source before every paid run.
- E

Paid setup traffic is crash-visible evidence. Before preflight, the command claims a sibling *--setup-traffic artifact and durably appends each completed row. A refused benchmark seals that setup-only artifact with performance_result=false, sla_result=false, and capacity_result=false. If --force continues after an unreadable answer, the durable outcome is preflight_forced_unreadable, never passed, and every downstream run or sweep remains INVALID diagnostic evidence.

STAGE ONE, CONTINUED

Exercise the endpoint contract without making a timing claim.

This command deterministically plans **two preflight rows + one calibration row + one measured replay row**, with no probes. Its illustrative output-budget p50/p95 is **320/480 tokens**, deriving a **720-token request cap**. The default fallback envelope permits at most **12 physical POST attempts**. With the explicit managed thinking-off control, its offline peak plan is **89,202 input and 4,464 output tokens/minute**. Those are conservative admission reservations, not observed usage. It qualifies mechanics only: **no report-latency, throughput, SLA, or capacity claim**.

GLM 5.2 INSTRUMENT CANARY

REAL, BILLABLE REQUESTS

```
python3 -m traffic_replay benchmark \
  --host https://YOUR-WORKSPACE-HOST \
  --endpoint databricks-glm-5-2 \
  --auth-profile YOUR-DATABRICKS-PROFILE \
  --profile configs/profile_glm52_canary_illustrative.json \
  --extra-body '{"reasoning_effort":"none"}' \
  --fixed-rate 0.1 --duration 12 \
  --ttft-definition first_visible \
  --rate-limits configs/rate_limits_databricks_glm_5_2_enterprise_p2t_2026-08-07.json \
  --out-dir results/glm52-instrument-canary \
  --label "INSTRUMENT CONFORMANCE ONLY – NOT CUSTOMER DEMAND OR CAPACITY"
```

1. Exact representatives

Profile mode sends its concrete p50 and p95 requests. Prompt mode sends the first two approved prompts, cycling only if one exists.

2. Exact pass condition

HTTP 200 plus non-refusal visible content or a valid non-refusal tool call, clean stream completion, and zero parse errors.

3. Exact refusal

If either representative is incomplete, the measured load stops unless the operator explicitly uses `--force`. Force permits only an explicitly `INVALID` diagnostic after reachable HTTP responses; it does not turn the gate into a pass or override transport failure.

R Managed GLM 5.2 reasoning control: Owner-confirmed behavior establishes that top-level `{"reasoning_effort":"none"}` disables reasoning on Unity AI Gateway `system.ai.glm-5-2` and the direct endpoint; omission selects maximum reasoning. The public guide names the field without enumerating GLM values. This release qualifies only direct `/serving-endpoints/.../invocations`; Gateway is protocol-diagnostic because destination identity, routing/fallback, and combined quotas are not bound, so it supports no quota or capacity conclusion. Direct SGLang uses `{"chat_template_kwargs":{"enable_thinking":false}}`; with thinking on, unset is Max and "high" is High. The canary selects no reasoning and must still pass full preflight; HTTP 200 alone is insufficient.

[Databricks limits and quotas](#) [Databricks reasoning-model guide](#) [Unity AI Gateway model services](#)

STAGE TWO

Run the customer contract, then separate SLA from capacity.

Build configs/profile_measured.json from approved request metrics, review its digest and dropped rows, then set rates, customer-owned targets, and client bounds from the authorized experiment.

MEASURED FIXED-RATE RUN

TEMPLATE - REPLACE EVERY YOUR_* VALUE

```
python3 -m traffic_replay benchmark \
  --host https://YOUR-WORKSPACE-HOST \
  --endpoint YOUR-ENDPOINT-NAME \
  --auth-profile YOUR-DATABRICKS-PROFILE \
  --profile configs/profile_measured.json \
  --fixed-rate YOUR-AUTHORIZED-RPS --duration 300 \
  --max-concurrency YOUR-TESTED-WORKER-BOUND \
  --max-pending-requests YOUR-TESTED-PENDING-BOUND \
  --ttft-definition first_visible \
  --ttft-p95 YOUR-TTFT-MS --ttfg-p95 YOUR-TTFG-MS \
  --success-rate YOUR-FRACTION-STRICTLY-BETWEEN-0-AND-1 --rate-limits RATE_LIMITS.json \
  --fail-on caution --out-dir results/customer-fixed-rate
```

R **Arrival rate is the load axis.** Do not use legacy `--concurrency`. `--max-concurrency` is only a client safety bound; the report shows observed in-flight concurrency. The harness opens a fresh HTTP/1.1 connection per physical attempt, so capacity stays inconclusive for an unknown, pooled, or HTTP/2 production client. Add `--production-connection-policy fresh_http1_per_physical_attempt` only when that is the real application's exact policy; the recorded assertion is not an observation of production. For reasoning models this template explicitly scores `first_visible`; use `first_content` only when visible, reasoning, or refusal onset is the customer-owned target.

GUARDED RATE SWEEP

SEQUENTIAL AND STATEFUL

```
python3 -m traffic_replay sweep \
  --host https://YOUR-WORKSPACE-HOST --endpoint YOUR-ENDPOINT-NAME \
  --auth-profile YOUR-DATABRICKS-PROFILE \
  --profile configs/profile_measured.json \
  --rate YOUR-RUNG-1,YOUR-RUNG-2,YOUR-RUNG-3 \
  --duration 120 --cooldown 60 \
  --cpt YOUR-PREMEASURED-CHARACTERS-PER-TOKEN \
  --max-concurrency YOUR-TESTED-WORKER-BOUND \
  --max-pending-requests YOUR-TESTED-PENDING-BOUND \
  --ttft-definition first_visible \
  --ttft-p95 YOUR-TTFT-MS --ttfg-p95 YOUR-TTFG-MS \
  --success-rate YOUR-FRACTION-STRICTLY-BETWEEN-0-AND-1 --rate-limits RATE_LIMITS.json \
  --out-dir results/customer-rate-sweep
```

! Cooldown is spacing, not proof that rolling quotas, autoscaling, cache state, or unrelated traffic reset. The whole requested ladder is budgeted before preflight, then every physical POST is admitted again at runtime. Stop on local admission denial, 429s, parse errors, answer failures, target misses, load shortfall, pending drops, generator saturation, or production impact. A top rung that holds establishes no ceiling; extend only in a newly authorized window after quota, cost, generator, and endpoint telemetry review.

i **Interchunk is a stream-gap diagnostic, not token-to-token latency.** It is the maximum gap between consecutive nonempty visible, reasoning, or refusal SSE delta events. Tool fragments, usage-only events, heartbeats, and streams with only one eligible event are excluded. TTFB is the first nonempty bounded response-body chunk returned by the client read, not the first socket byte or first parsed SSE line.

STAGE THREE

Read independent gates before quoting any number.

A target pass can coexist with invalid identity, unstable deployment metadata, local quota denial, or weak measurement evidence. Preserve the exact qualification and the source directory.

Read five decisions, not one green badge

Identity, stability, and runtime admission are evidence gates within these five decisions, not extra canonical dimensions.

DIMENSION	QUESTION ANSWERED	REQUIRED READING
Evidence integrity	Do current bytes match the manifest and completion chain?	Verify externally; a fresh run cannot verify the manifest that encloses it.
Measurement validity	Did errors, delivery, parsing, sample size, response identity, transport mismatch, or endpoint stability undermine inference?	Mixed response models, mismatch with an explicit request model, an unexpected Databricks served-model-name, or changed endpoint metadata invalidates the claim; incomplete or production-mismatched transport evidence is caution.
Acceptance checks	Did the run meet the customer-owned explicit targets?	PASS, MISS, INCONCLUSIVE, or NOT EVALUATED. A reliability pass also requires the one-sided 95% Wilson lower confidence bound to meet the target.
Quota state	Was every physical POST locally admitted, and was terminal HTTP 429 observed on any supplied request row?	A local denial sends no POST and is not provider rate limiting. Zero observed 429s is not quota headroom. One row can contain multiple attempts.
Endpoint capacity	Did the endpoint hold at this exact tested load?	Never translate one held rung into a ceiling or guarantee.

Create a separate run-verification receipt

The verifier re-reads the sealed source, cross-checks the canonical journal and summary, and writes a new receipt without modifying the run.

```
RUN_DIR=results/customer-fixed-rate/RUN-DIRECTORY
python3 -m traffic_replay verify-run "$RUN_DIR" \
  --out "${RUN_DIR}-verification"

python3 -m traffic_replay verify-sweep \
  results/customer-rate-sweep
```

The receipt path resolves to results/customer-fixed-rate/RUN-DIRECTORY-verification.

Receipt boundary

verification.json, verified-report.html, and verified-report.md are sealed in the receipt. SHA-256 consistency is not a digital signature: it does not prove authorship, trusted time, repository availability, or immutability after verification.

Before sharing: review labels, endpoint metadata, file paths, and configuration. Run artifacts are designed to retain metadata and hashes rather than raw request or response content; the measured extra_body remains part of reproducibility evidence.

[Repository landing page](#)

Before external sharing, confirm the stamped source commit is published; the repository's default branch may not yet contain this revision.

Source commit 1cf342627397cfd246a7cefc3a551f1e5b020fc3 · canonical HTML SHA-256
45705a98efb8347d305573f952ecb691e0bc56be9da99abdcf3de1fdaa748d3d · source commit time 2026-08-
10T14:55:23Z